

# **Training Neural Networks**

## **Part 2**

Hongzhi Liu

November 22nd, 2019

# CONTENT

---



**1. Optimization**

**2. Learning Rate Decay**

**3. Regularization**

---



## 1. Gradient Descent (GD)

A one-dimensional Taylor series is an expansion of a real function about a point  $x=a$  is given by

$$f(x) = \sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} (x - a)^n$$

with Peano remainder

$$f(x + \epsilon) \approx f(x) + f'(x)\epsilon + \mathcal{O}(\epsilon^2)$$

$$f(x + \epsilon) \approx f(x) + f'(x)\epsilon$$

$$f(x - \eta f'(x)) \approx f(x) - \eta f'(x)^2 \quad \eta > 0$$

$$x \leftarrow x - \eta f'(x)$$



## 2. Stochastic Gradient Descent (SGD)

### SGD

```
while True:
    dx = compute_gradient(x)
    x -= learning_rate * dx
```

### SGD+Momentum

```
vx = 0
while True:
    dx = compute_gradient(x)
    vx = rho * vx + dx
    x -= learning_rate * vx
```

### SGD + Nesterov Momentum

```
v = 0
while True:
    dx = compute_gradient(x)
    old_v = v
    v = rho * v - learning_rate * dx
    x += -rho * old_v + (1 + rho) * v
```



## 3. AdaGrad and RMSProp

### AdaGrad

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared += dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7))
```

### RMSProp: “Leaky AdaGrad”

```
grad_squared = 0
while True:
    dx = compute_gradient(x)
    grad_squared += decay_rate * grad_squared + (1 - decay_rate) * dx * dx
    x -= learning_rate * dx / (np.sqrt(grad_squared) + 1e-7))
```



## 4. Adam

### Full form

```
first_moment = 0
second_moment = 0
for t in range(1, num_iterations):
    dx = compute_gradient(x)
    first_moment = beta1 * first_moment + (1 - beta1) * dx
    second_moment = beta2 * second_moment + (1 - beta2) * dx * dx
    first_unbias = first_moment / (1 - beta1 ** t)
    second_unbias = second_moment / (1 - beta2 ** t)
    grad_squared += decay_rate * grad_squared + (1 - decay_rate) * dx * dx
    x -= learning_rate * first_unbias / (np.sqrt(second_unbias) + 1e-7))
```



# Learning Rate Decay



Linear  $\alpha_t = \alpha_0(1 - t/T)$

```
for t in range(1, epoch):  
    learning_rate = learning_rate * (1 - t / epoch)
```

Inverse sqrt  $\alpha_t = \alpha_0/\sqrt{t}$

```
for t in range(1, epoch):  
    learning_rate = learning_rate / (sqrt(t))
```

Cosine  $\alpha_t = \frac{1}{2}\alpha_0(1 + \cos(t\pi/T))$

```
for t in range(1, epoch):  
    learning_rate = 0.5 * learning_rate * (1 + cos(pi * t / epoch))
```



## 1. Dropout

```
p = ratio
def train_step(X):
    H_1 = np.maximum(0, np.dot(W1, X) + b_1)
    U_1 = np.random.rand(*H_1.shape) < p
    H_1 *= U_1
    H_2 = np.maximum(0, np.dot(W2, X) + b_2)
    U_2 = np.random.rand(*H_2.shape) < p
    H_2 *= U_2
    out = np.dot(W3, H2) + b3
```





## 2. Batch Normalization (BN)

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

1. It accelerates the training of deep neural nets.

2. It also acts as a regularizer, in some cases eliminating the need for Dropout.

3. It is helpful to reduce the over-fitting of the network and improve the accuracy rate as a whole.



## 3. Data Augmentation

### Horizontal Flips

```
class torchvision.transforms.RandomHorizontalFlip(p=0.5)
```

### Random crops

```
class torchvision.transforms.RandomCrop(size, padding=None,  
                                         pad_if_needed=False, fill=0, padding_mode='constant')
```

### Scales

```
class torchvision.transforms.Scale(size, interpolation=2)
```

### ColorJitter

```
class torchvision.transforms.ColorJitter(brightness=0, contrast=0,  
                                         saturation=0, hue=0)
```

**Thanks**